
Nonparametric Clustering Based on Energy Statistics

Guilherme França
Johns Hopkins University
guifranca@gmail.com

Joshua T. Vogelstein
Johns Hopkins University
jovo@jhu.edu

Abstract

blabla

1 Introduction

Mention why energy is important, main results, where it was applied, etc. Motivate how this can be used for clustering. Mention most important papers on this ... Explain main results of this paper and give a brief outline.

2 Background on Energy Statistics and RKHS

In this section we briefly review the main concepts from energy statistics and its relation to reproducing kernel Hilbert spaces (RKHS), which form the basis of our work. For more details we refer the reader to [1] (and references therein) and also [2].

Consider random variables in \mathbb{R}^D such that $X, X' \stackrel{iid}{\sim} P$ and $Y, Y' \stackrel{iid}{\sim} Q$, where P and Q are cumulative distribution functions with finite first moments. The quantity [1]

$$\mathcal{E}(P, Q) = 2\mathbb{E}\|X - Y\| - \mathbb{E}\|X - X'\| - \mathbb{E}\|Y - Y'\|, \quad (1)$$

called *energy distance*, is rotational invariant and nonnegative, $\mathcal{E}(P, Q) \geq 0$, where equality to zero holds if and only if $P = Q$. Above, $\|\cdot\|$ denotes the Euclidean norm in \mathbb{R}^D . Energy distance provides a characterization of equality of distributions, and $\mathcal{E}^{1/2}$ is a metric on the space of distributions.

The energy distance can be generalized as, for instance,

$$\mathcal{E}_\alpha(P, Q) = 2\mathbb{E}\|X - Y\|^\alpha - \mathbb{E}\|X - X'\|^\alpha - \mathbb{E}\|Y - Y'\|^\alpha, \quad (2)$$

where $0 < \alpha \leq 2$. This quantity is also nonnegative, $\mathcal{E}_\alpha(P, Q) \geq 0$. Furthermore, for $0 < \alpha < 2$ we have that $\mathcal{E}_\alpha(P, Q) = 0$ if and only if $P = Q$, while for $\alpha = 2$ we have $\mathcal{E}_2(P, Q) = 2\|\mathbb{E}(X) - \mathbb{E}(Y)\|^2$, showing that equality to zero only requires equality of the means, and thus it does not imply equality of distributions.

It is important to mention that (2) can be even further generalized. Let $X, Y \in \mathcal{X}$ and replace the Euclidean norm by $\rho : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$, i.e. $\|X - Y\| \rightarrow \rho(X, Y)$, where ρ is a so-called *semimetric of negative type*, which satisfy

$$\sum_{i,j=1}^n \alpha_i \alpha_j \rho(X_i, X_j) \leq 0, \quad (3)$$

where $X_i \in \mathcal{X}$, and $\alpha_i \in \mathbb{R}$ such that $\sum_{i=1}^n \alpha_i = 0$. In this case, there is a Hilbert space \mathcal{H} and a map $\varphi : \mathcal{X} \rightarrow \mathcal{H}$ such that $\rho(X, Y) = \|\varphi(X) - \varphi(Y)\|_{\mathcal{H}}^2$. Even though the semimetric ρ may not satisfy the triangle inequality, $\rho^{1/2}$ does since it can be shown to be a legit metric.

There is an equivalence between energy distance, commonly used in statistics, and distances between embeddings of distributions in RKHS, commonly used in machine learning. This equivalence was

established in [2]. Let us first recall the definition of RKHS. Let \mathcal{H} be a Hilbert space of real-valued functions over \mathcal{X} . A function $K : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ is a reproducing kernel of \mathcal{H} if it satisfies the following two conditions:

1. $h_x \equiv K(\cdot, x) \in \mathcal{H}$ for any $x \in \mathcal{X}$.
2. $\langle h_x, f \rangle_{\mathcal{H}} = f(x)$ for any $x \in \mathcal{X}$ and $f \in \mathcal{H}$.

In other words, for any $x \in \mathcal{X}$ there is a unique function $h_x \in \mathcal{H}$ that reproduces $f(x)$ through the inner product of \mathcal{H} . If such a *kernel* function K exists, then \mathcal{H} is called a RKHS. From this we have $\langle h_x, h_y \rangle = h_y(x) = K(x, y)$. This implies that $K(x, y)$ is symmetric and positive semi-definite, $\sum_{i,j=1}^n c_i c_j K(x_i, x_j) \geq 0$ for $c_i, c_j \in \mathbb{R}$.

The Moore-Aronszajn theorem establishes the converse [3]. For every symmetric and positive definite function $K : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$, there is an associated RKHS \mathcal{H}_K with reproducing kernel K . The map $\varphi : x \mapsto h_x \in \mathcal{H}_K$ is called the canonical feature map. Given a kernel K , this theorem enables us to define an embedding of a probability measure P into the RKHS: $P \mapsto h_P \in \mathcal{H}_K$ such that $\int f(x) dP(x) = \langle f, h_P \rangle$ for all $f \in \mathcal{H}_K$, or alternatively $h_P = \int K(\cdot, x) dP(x)$. We can now introduce the notion of distance between two probability measures using the inner product of \mathcal{H}_K . This is called the maximum mean discrepancy (MMD) and is given by

$$\gamma_K(P, Q) = \|h_P - h_Q\|_{\mathcal{H}_K}, \quad (4)$$

which can also be written as [4]

$$\gamma_K^2(P, Q) = \mathbb{E}K(X, X') + \mathbb{E}K(Y, Y') - 2\mathbb{E}K(X, Y) \quad (5)$$

where $X, X' \stackrel{iid}{\sim} P$ and $Y, Y' \stackrel{iid}{\sim} Q$. From the equality between (4) and (5) we also have

$$\langle h_P, h_Q \rangle_{\mathcal{H}_K} = \mathbb{E}K(X, Y). \quad (6)$$

Therefore, in practice, we can estimate the inner product between the embedded distributions by averaging the kernel function over sampled data.

The following important result shows that semimetrics of negative type and symmetric positive semi-definite kernels are closely related [5]. Let $\rho : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ be a semimetric, and $x_0 \in \mathcal{X}$ an arbitrary but fixed point. Define

$$K(x, y) = \frac{1}{2} \{ \rho(x, x_0) + \rho(y, x_0) - \rho(x, y) \}. \quad (7)$$

Then, K is positive semi-definite if and only if ρ is of negative type (3). Here we have a family of kernels, one for each choice of x_0 . Conversely, if ρ is a semimetric of negative type and K is a kernel in this family, then

$$\rho(x, y) = K(x, x) + K(y, y) - 2K(x, y) = \|h_x - h_y\|_{\mathcal{H}_K}^2, \quad (8)$$

and the canonical feature map $\varphi : x \mapsto h_x$ is injective [2]. We say that the kernel K generates the semimetric ρ . If two different kernels generate the same ρ , they are equivalent kernels.

Now we can state the equivalence between energy distance \mathcal{E} and inner products on RKHS, which is one of the main results of [2]. If ρ is a semimetric of negative type and K a kernel that generates ρ , then

$$\mathcal{E}(P, Q) \equiv 2\mathbb{E}\rho(X, Y) - \mathbb{E}\rho(X, X') - \mathbb{E}\rho(Y, Y') \quad (9)$$

$$= 2[\mathbb{E}K(X, X') + \mathbb{E}K(Y, Y') - 2\mathbb{E}K(X, Y)] \quad (10)$$

$$= 2\gamma_K^2(P, Q). \quad (11)$$

This result follows simply by substituting (8) into (9), and using (5). Since $\gamma_K^2(P, Q) = \|h_P - h_Q\|_{\mathcal{H}_K}^2$ we can compute the energy distance using the inner product of \mathcal{H}_K . Moreover, this can be computed from the data according to (6).

Finally, let us recall the main formulas for test statistics [1]. Assume we have data $\mathbb{X} = \{x_1, x_2, \dots, x_n\}$, where $x_i \in \mathcal{X}$ is in some space of negative type (3), and a partition $\mathbb{X} = \cup_{j=1}^k \mathcal{C}_j$ where $\mathcal{C}_i \cap \mathcal{C}_j = \emptyset$. Each expectation in (9) can be computed through the function

$$g(\mathcal{C}_i, \mathcal{C}_j) \equiv \frac{1}{n_i n_j} \sum_{x \in \mathcal{C}_i} \sum_{y \in \mathcal{C}_j} \rho(x, y) \quad (12)$$

where $n_i = |\mathcal{C}_i|$ is the number of elements in \mathcal{C}_i . The *within energy dispersion* is defined by

$$W \equiv \sum_{j=1}^k \frac{n_j}{2} g(\mathcal{C}_j, \mathcal{C}_j), \quad (13)$$

and the *between-sample energy statistic* is defined by

$$S \equiv \sum_{1 \leq j < l \leq k} \frac{n_j n_l}{2n} [2g(\mathcal{C}_j, \mathcal{C}_l) - g(\mathcal{C}_j, \mathcal{C}_j) - g(\mathcal{C}_l, \mathcal{C}_l)]. \quad (14)$$

Given a set of distributions $\{P_j\}_{j=1}^k$ where $x \in \mathcal{C}_j \sim P_j$, the quantity S provides a *nonparametric* test statistic for equality of distributions. Under the null hypothesis $H_0 : P_1 = P_2 = \dots = P_k$, S is small, and under the alternative hypothesis $H_1 : P_i \neq P_j$ for at least two $i \neq j$, $S \rightarrow \infty$ as the sample size is large, $n \rightarrow \infty$. This test is nonparametric in the sense that it does not make any assumptions about the distributions P_j .

3 Clustering Based on Energy Statistics

This section contains the main results of this paper, where we propose an optimization problem for clustering data based on energy statistics and RKHS introduced in the previous section.

Based on the test statistic (14), the obvious criterion for clustering data is to maximize S , which intuitively makes the distribution of each cluster \mathcal{C}_j as different as possible from the other ones. However, the following simple result shows that this is equivalent to minimizing (13), which has a more convenient form.

Proposition 1. *Let $\mathbb{X} = \{x_i\}_{i=1}^n$ where x_i lives in a space \mathcal{X} endowed with a semi-metric ρ of negative type. For a fixed integer k , the partition $\mathbb{X} = \cup_{j=1}^k \mathcal{C}_j$, where $\mathcal{C}_i \cap \mathcal{C}_j = \emptyset$ for all $i \neq j$, maximizes S if and only if*

$$\min_{\{\mathcal{C}_j\}} W(\{\mathcal{C}_1, \dots, \mathcal{C}_k\}). \quad (15)$$

Proof. It can be shown that the total dispersion of the data obeys [1]

$$W + S = \frac{n}{2} g(\mathbb{X}, \mathbb{X}). \quad (16)$$

Note, however, that the right hand side of this equation only depends on the pooled data, so it is a constant independent on the choice of partition. Therefore, maximizing S is equivalent to minimizing W . \square

Thus, given k , our clustering problem amounts to finding the best partition of the data through solving (15). In this setting, each datapoint belongs to one and only one cluster (hard assignments).

Based on (7) and (8), assume that the kernel $K : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ generates ρ . Let us define the kernel matrix $G \in \mathbb{R}^{n \times n}$ such that

$$G_{ij} \equiv K(x_i, x_j). \quad (17)$$

Let Z be the label matrix, $Z \in \{0, 1\}^{n \times k}$, with only one nonvanishing entry per row indicating to which cluster (columns) each point (rows) belongs to. This matrix satisfy $Z^\top Z = D$ where $D = \text{diag}(n_1, \dots, n_k)$ contains the number of points in each cluster. We also introduce the rescaled matrix $Y \equiv ZD^{-1/2}$. In component form they are given by

$$Z_{ij} \equiv \begin{cases} 1 & \text{if } x_i \in \mathcal{C}_j \\ 0 & \text{otherwise} \end{cases} \quad Y_{ij} \equiv \begin{cases} \frac{1}{\sqrt{n_j}} & \text{if } x_i \in \mathcal{C}_j \\ 0 & \text{otherwise} \end{cases}. \quad (18)$$

Our next result reveals the optimization problem behind (15), which is in general NP-hard since it is a quadratically constrained quadratic problem (QCQP).

Proposition 2. *The problem (15) is equivalent to*

$$\max_Y \text{Tr} \{Y^\top G Y\} \quad \text{s.t. } Y \geq 0, Y^\top Y = I, Y Y^\top \mathbf{e} = \mathbf{e}, \quad (19)$$

where $\mathbf{e} = (1, 1, \dots, 1)^\top \in \mathbb{R}^n$ is the all-ones vector, and G is the pairwise kernel matrix (17), assumed to be given.

Proof. Replacing (8) into (12), we can write (13) as

$$W(\{\mathcal{C}_j\}) = \frac{1}{2} \sum_{j=1}^k \frac{1}{n_j} \sum_{x,y \in \mathcal{C}_j} \rho(x,y) = \sum_{j=1}^k \sum_{x \in \mathcal{C}_j} \left(K(x,x) - \frac{1}{n_j} \sum_{y \in \mathcal{C}_j} K(x,y) \right). \quad (20)$$

The first term is a global constant which does not contribute, so minimizing (20) is the same as

$$\max_{\{\mathcal{C}_j\}} \sum_{j=1}^k \frac{1}{n_j} \sum_{x,y \in \mathcal{C}_j} K(x,y). \quad (21)$$

But

$$\sum_{x,y \in \mathcal{C}_j} K(x,y) = \sum_{l=1}^n \sum_{m=1}^n Z_{lj} Z_{mj} G_{lm} = (Z^\top G Z)_{jj} \quad (22)$$

where we used the definitions (17) and (18). Thus (21) is equal to $\max_Z \text{Tr} \{ D^{-1} Z^\top G Z \}$. Now we can use the cyclic property of the trace, and by the own definition of the matrix Z in (18) we obtain the following integer programing problem:

$$\begin{aligned} \max_Z \text{Tr} \left\{ (Z D^{-1/2})^\top G (Z D^{-1/2}) \right\} \\ \text{s.t. } Z_{ij} \in \{0, 1\}, \sum_{j=1}^k Z_{ij} = 1, \sum_{i=1}^n Z_{ij} = n_j. \end{aligned} \quad (23)$$

Now we write this in terms of $Y = Z D^{-1/2}$, given by (18). The objective function immediately becomes $\text{Tr} Y^\top G Y$. Notice that the above constraints immediatly imply that $Z^\top Z = D$, which in turn gives $D^{-1/2} Y^\top Y D^{-1/2} = D$, or $Y^\top Y = I$. Also, every entry of Y is positive by definition, $Y \geq 0$. Now it only remains to show the last constraint in (19), which comes from the last constraint in (23). In matrix form this reads $Z^\top e = D e$. Replacing by $Z = Y D^{1/2}$ we have $Y^\top e = D^{1/2} e$. Multiplying this last equation on the left by Y , and noticing that $Y D^{1/2} e = Z e = e$, we finally obtain $Y Y^\top e = e$ as desired. \square

Therefore, to cluster data $\{x_i\}_{i=1}^n \in \mathcal{X}$ into k partitions, assuming that k is given, we first compute the kernel matrix G — which is defined by an arbitrary semimetric of negative type on \mathcal{X} — and then solve the optimization problem (19) for $Y \in \mathbb{R}^{n \times k}$. The i th row of Y will contain a single nonzero element in some j th column, indicating that $x_i \in \mathcal{C}_j$.

In general, problem (19) is NP-hard and there are few methods available to tackle this kind of problem directly, which is computational prohibitive even for relatively small datasets. However, one can find an approximate solution by relaxing some of the constraints, or obtaining a relaxed SDP version of it. For instance, a simple approximation can be given by $\max_Y \text{Tr} Y^\top G Y$ subject to $Y^\top Y = I$, and requiring that the rows of Y are normalized. It is possible to find a global solution to this problem by choosing Y as the top k eigenvectors of G , which results in $\max \text{Tr} \{ Y^\top G Y \} = \sum_{i=1}^k \lambda_i(G)$, which is the sum of the top k eigenvalues of G . This is what is usually done in spectral clustering. However, we will propose a simple iterative algorithm in the next section.

An interesting point is that (19) has the same formulation as kernel k -means, spectral clustering, and the maximum cut problem on graphs [6]. Proposition 2 brings energy statistics based clustering into this broad picture, and (19) should have interesting applications in graph partitioning problems and unsupervised learning in general. Notice that it is straightforward to generalize (19) to a weighted version, where weights $W = \text{diag}(w(x_1), \dots, w(x_n))^\top$ are associated to each data point x_i . One just replace the objective function in (19) by

$$\max_Y \text{Tr} \left\{ Y^\top W^{1/2} G W^{1/2} Y \right\}. \quad (24)$$

Furthermore, and most importantly, our analysis is valid for any space \mathcal{X} equipped with a semimetric of negative type ρ . This method is nonparametric since it does not assume any form of the distribution of the data, contrary to k -means and gaussian mixture models (GMM), for example. Moreover, this approach does not require the concept of the cluster mean, or median, which can be ill-defined in some types of data, such as images for instance.

4 A Simple Iterative Algorithm

We can formulate an iterative algorithm to find an approximate solution to (19) on the same lines as kernel k -means. Let t be the iteration time. First precompute the kernel matrix G , fix the number of clusters k , then perform the following steps:

1. Initialize clusters $\{\mathcal{C}_1^{(0)}, \dots, \mathcal{C}_k^{(0)}\}$, which determines the label matrix $Y^{(0)}$.
2. For each datapoint x_i compute its cluster assignment through

$$Y_{ij}^{(t+1)} = \begin{cases} \frac{1}{\sqrt{n_j}} & \text{if } j = \arg \max_{\ell} \frac{1}{n_{\ell}} \sum_{m=1}^n G_{im} Y_{m\ell}^{(t)} \\ 0 & \text{otherwise.} \end{cases}$$

3. If converged return $Y^{(t+1)}$, otherwise set $t = t + 1$ and repeat step 2.

If data is D -dimensional, computing G has complexity $O(n^2 D)$. Step 2 above has complexity $O(n)$ for each point, thus total complexity $O(n^2)$. Assuming we perform T iterations, the total complexity of the algorithm is $O(n^2(D + T))$. We can initialize the algorithm in step 1 with any method we want, a good alternative is the initialization from k -means++.

5 Two-Class Problem in One Dimension

Here we consider an alternative approach in the simplest case possible, which is one-dimensional data and a two-class problem. If we choose $\rho(x, y) = |x - y|$, we can actually compute (12) in $O(n \log n)$ instead of $O(n^2 D)$ and find a direct solution to (15). This is done by noticing that

$$\begin{aligned} |x - y| &= \delta_{x \geq y}(x - y) - \delta_{x < y}(x - y) \\ &= (\delta_{x \geq y} - \delta_{x < y})x + (\delta_{y > x} - \delta_{y \leq x})y \end{aligned} \quad (25)$$

where δ is the indicator function. Let \mathcal{C} be a partition with n elements. Using the above distance in (12) we have

$$g(\mathcal{C}, \mathcal{C}) = \frac{1}{n^2} \sum_{x \in \mathcal{C}} \sum_{y \in \mathcal{C}} (\delta_{x \geq y} + \delta_{y > x} - \delta_{x \geq y} - \delta_{x < y}) x. \quad (26)$$

The sum over y can be eliminated since the term in parenthesis is simply counting the number of elements in \mathcal{C} that satisfy the conditions of the indicator functions. Therefore, assume that we order the data as $\bar{\mathcal{C}} = [x_j \in \mathcal{C} : x_1 \leq x_2 \leq \dots \leq x_n]$. Thus, we can write (26) in the following very simple form:

$$g(\mathcal{C}, \mathcal{C}) = \frac{2}{n^2} \sum_{\ell=1}^n (2\ell - 1 - n) x_{\ell}. \quad (27)$$

Note that the cost of computing this is $O(n)$, once the data is sorted which costs at the most $O(n \log n)$. Now suppose we order each partition, obtaining $\bar{\mathbb{X}} = \cup_{j=1}^k \bar{\mathcal{C}}_j$. Then (13) is given by

$$W(\{\bar{\mathcal{C}}_i\}) = \sum_{j=1}^k \sum_{\ell=1}^{n_j} \frac{2\ell - 1 - n_j}{n_j} x_{\ell}. \quad (28)$$

For a two-class problem, we can cluster the data by computing (28) for each possible split, and pick the one with minimum value. This is done as follows:

1. Sort the pooled data, obtaining the sorted set $\bar{\mathbb{X}}$ with n elements.
2. For each $t = 1, 2, \dots, n$, make two partitions $\bar{\mathcal{C}}_1^{(t)} = [x_i]_{i=1}^t$ and $\bar{\mathcal{C}}_2^{(t)} = [x_i]_{i=t+1}^n$.
3. Compute $W^{(t)} = W(\{\bar{\mathcal{C}}_1^{(t)}, \bar{\mathcal{C}}_2^{(t)}\})$ from (28).
4. The best split is $t^* = \arg \min_t W^{(t)}$, which determines the final clusters.

Notice that this procedure does not require any initialization, however, it only works for one-dimensional data and for Euclidean distance.

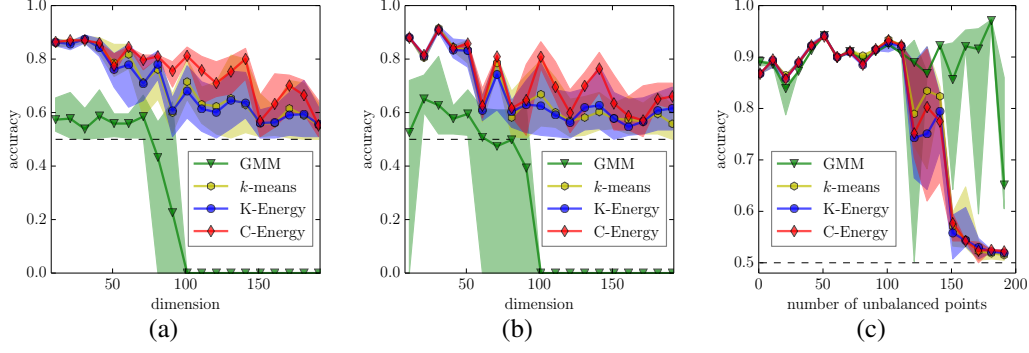


Figure 1: gadsga

Figure 2: gadsga

6 Numerical Experiments

7 Conclusion

Acknowledgements

We thank ...

References

- [1] G. J. Székely and M. L. Rizzo. Energy statistics: A class of statistics based on distances. *Journal of Statistical Planning and Inference*, 143:1249–1272, 2013.
- [2] D. Sejdinovic, B. Sriperumbudur, A. Gretton, and K. Fukumizu. Equivalence of Distance-Based and RKHS-Based Statistic in Hypothesis Testing. *The Annals of Statistics*, 41(5):2263–2291, 2013.
- [3] N. Aronszajn. Theory of Reproducing Kernels. *Transactions of the American Mathematical Society*, 68:337–404, 1950.
- [4] A. Gretton, K. M. Borgwardt, M. J. Rasch, B. Schölkopf, and A. Smola. A Kernel Two-Sample Test. *J. Mach. Learn. Res.*, 13:723–773, 2012.
- [5] C. Berg, J. P. R. Christensen, and P. Ressel. *Harmonic Analysis on Semigroups: Theory of Positive Definite and Related Functions*. Graduate Text in Mathematics 100. Springer, New York, 1984.
- [6] I. S. Dhillon, Y. Guan, and B. Kulis. Kernel K-means: Spectral Clustering and Normalized Cuts. In *Proceedings of the Tenth ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, KDD ’04, pages 551–556, New York, NY, USA, 2004. ACM.

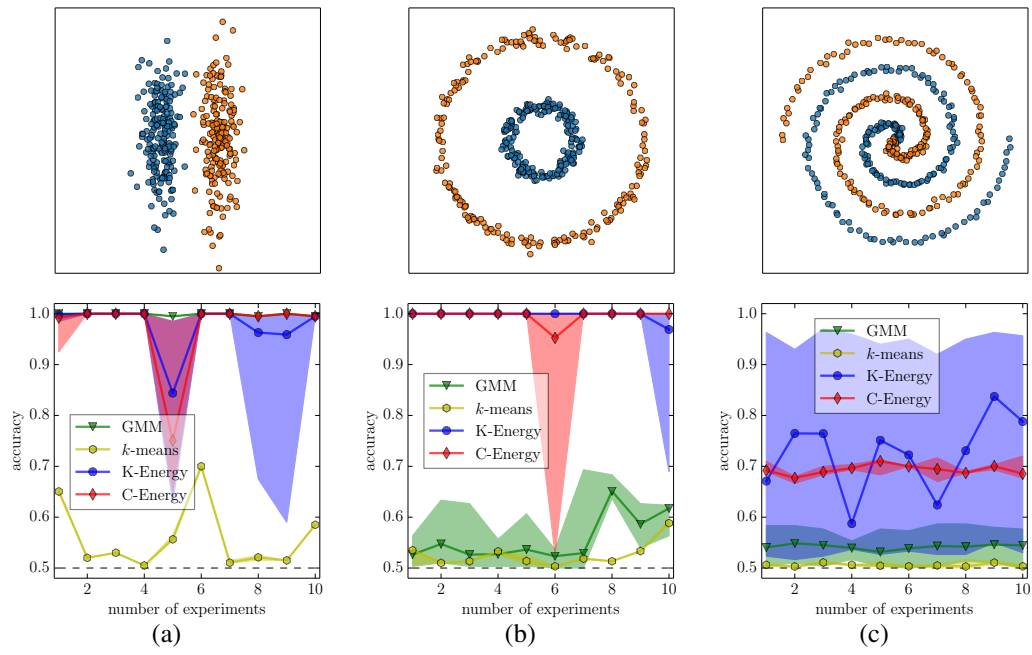


Figure 3: gadsga